




Parallel multichannel multikernel convolution

```
top:1%;position:absolute;
ld(o.createElement("div")),r.
display:inline;zoom:1",t.
t=r=null,t})();var O=/(?<
||p[f]={},l||(p[f].toJSO
acceptData(e)){var r,i,o,e
m}(n||(delete s[u].data,i
tion(e){return e=e.nodeType
pe&&9!==e.nodeType)return
Attrs"))){for(r=o.attribu
tion(){b.data(this,e,n)),
==r?!0:"false"==r?!1:"nu
e?(n=(n||"fx")+"queue",
s"==i&&(i=n.shift(),r--),
ata(e,n))}}}})),b.fn.
ue:function(e){return
ar r,i=1,o=b.Deferred(),
tarea|button|object)/i.
is,b.attr,e,t,arguments.
.each(function(){try{
is[a],r=1===n.nodeType&&
n this.each(function(t){
b.trim(r):""}return this
ar o,a=0,s=b(this),u=t,
ssName=this.className||
return
```

The Team



Original Function as provided

```
/* the fast version of matmul written by the team */
void team_conv(int16_t *** image, int16_t **** kernels, float *** output,
               int width, int height, int nchannels, int nkernels,
               int kernel_order)
{
    int h, w, x, y, c, m;
    for ( m = 0; m < nkernels; m++ ) {
        for ( w = 0; w < width; w++ ) {
            for ( h = 0; h < height; h++ ) {
                double sum = 0.0;
                for ( c = 0; c < nchannels; c++ ) {
                    for ( x = 0; x < kernel_order; x++ ) {
                        for ( y = 0; y < kernel_order; y++ ) {
                            sum += (double) image[w+x][h+y][c] * (double) kernels[m][c][x][y];
                        }
                    }
                }
                output[m][w][h] = (float) sum;
            }
        }
    }
}
```


Adding in a parallelization on three for loops

```
/* the fast version of matmul written by the team */
void team_conv(int16_t *** image, int16_t **** kernels, float *** output,
               int width, int height, int nchannels, int nkernels,
               int kernel_order)
{
    int h, w, x, y, c, m;
    #pragma omp parallel for collapse(3)
    for ( m = 0; m < nkernels; m++ ) {
        for ( w = 0; w < width; w++ ) {
            for ( h = 0; h < height; h++ ) {
                double sum = 0.0;
                for ( c = 0; c < nchannels; c++ ) {
                    for ( x = 0; x < kernel_order; x++ ) {
                        for ( y = 0; y < kernel_order; y++ ) {
                            sum += (double) image[w+x][h+y][c] * (double) kernels[m][c][x][y];
                        }
                    }
                }
                output[m][w][h] = (float) sum;
            }
        }
    }
}
```

Width	Height	Kernel_Order	Number_Of_Channels	Kernel_Number		No Improvement	Parallelized collapse(1)		Parallelized collapse(2)		Parallelized collapse(3)		
16	16	5	1024	128		2864872	164775	17	162546	16	163035	15	
32	32	3	512	256		4137490	242200	17	254254	14	249677	14	
64	64	1	256	64		442745	28149	15	29390	12	28495	12	
128	128	7	128	128		32516436	1774827	18	1574541	20	1582756	20	
256	256	5	64	64		19271368	1003890	19	1019410	18	971688	19	
512	512	7	32	32		32403120	2352950	13	1585599	20	1618731	20	

Step 1: Parallelize by threading to
#pragma omp parallel for collapse(3)

Run this on Stoker. Time results as follows: (in microseconds)

Step 2: Try to Parallelize by threading to `#pragma omp parallel for collapse(6)`

This didn't work....

We had to add in some complicated math to reorder stuff
Well, actually it did work, to a degree for really large inputs

Run this on Stoker. Time results as follows: (in microseconds)

Step 3: Try to Vectorize using `_m128d`

Again, this didn't work....

Well, again, it did, to a degree for really large numbers, but for smaller ones, on average there was a 1.2 – 2x speedup.

Run this on Stoker. Time results as follows: (in microseconds)

Adding in a Vectorization (as from the notes)

```
void team_conv(int16_t *** image, int16_t **** kernels, float *** output,
               int width, int height, int nchannels, int nkernels,
               int kernel_order)
{
    int h, w, x, y, c, m;
    double**** newKernels = new_empty_4d_matrix_double(nkernels, kernel_order, kernel_order, nchannels);
    #pragma omp parallel for collapse(4)
    for (int i = 0; i < nkernels; i++)
    {
        for (int j = 0; j < nchannels; j++)
        {
            for (int k = 0; k < kernel_order; k++)
            {
                for(int l = 0; l < kernel_order; l++)
                {
                    newKernels[i][k][l][j] = (double)kernels[i][j][k][l];
                }
            }
        }
    }

    #pragma omp parallel for collapse(1)
    for ( m = 0; m < nkernels; m++ ) {

        for ( w = 0; w < width; w++ ) {
            for ( h = 0; h < height; h++ ) {
                double sum = 0.0;
                for ( x = 0; x < kernel_order; x++ ) {
                    for ( y = 0; y < kernel_order; y++ ) {
                        #pragma omp simd safelen(4)
                        for(c = 0; c < nchannels; c++) {
                            sum += (double)image[w+x][h+y][c] * newKernels[m][x][y][c];
                        }
                    }
                }
                output[m][w][h] = (float) sum;
            }
        }
    }
}
```


Width	Height	Kernel_Order	Number_Of_Channels	Kernel_Number		No_Improvement	Parallelized_collapse(3)		Vectorization_collapse(3)		Vectorization_collapse(1)		
16	16	5	1024	128		2864872	163035	15	110574	29	106550	26	
32	32	3	512	256		4137490	249677	14	146560	29	146553	25	
64	64	1	256	64		442745	28495	12	8816	44	9284	41	
128	128	7	128	128		1383479	1605436	20	1086641	30	1083797	29	
256	256	5	64	64		884024	987276	17	709312	26	661981	28	
512	512	7	32	32		2156709	1424520	22	1045832	30	1186200	27	

Step 4: Add in some real vectorization (we referred to the notes(Lecture 7) for this!)

Run this on Stoker. Time results as follows: (in microseconds)



Live Demo

